

TOXIC COMMENTS CLASSIFICATION

IN SOCIAL NETWORKING

Submitted By

PENTAPATI ESWAR AKHIL

39110765

INDEX

1. Introduction

- a. Overview
- b. Purpose

1. Literature Survey

- a. Existing Problem
- b. Proposed solution

2. Theoretical Analysis

- a. Block Diagram
- b. Hardware/Software designing

3. Experimental investigations

4. Flowchart

5. Result

7. Advantages & Disadvantages

1. Applications

2. Conclusion

3. Future Scope

4. Bibliography

12. Appendix

- a. Source Code
- b. UI output screen-shot

1. INTRODUCTION

a. Overview:

In recent years, the flow of data over the internet has grown dramatically, especially with the appearance of social networking sites. Social networks sometimes become a place for threats, insults, and other components of cyber bullying. A huge number of people are involved in online social networks.

Toxic comments are textual comments with threats, insults, obscene, racism, etc. In recent years there have been many cases in which authorities have arrested some users of social sites because of the negative (abusive) content of their pages. Hence, the protection of network users from anti-social behavior is an important activity. One of the major tasks of such activity is automated detecting the toxic comments.

The project aims to build a multi-headed model that's capable of detecting different types of toxicity like threats, obscenity, insults, and identity-based hate, we will use a dataset of comments from Wikipedia talk page edits, collected by Kaggle. Improvements to the current model

b. Purpose:

We'll be able to understand the problem to classify if it is a toxic comment or a normal comment. We will be able to know how to pre-process/clean the data using different data pre-processing techniques. You will be able to analyse or get insights into data through visualization. Applying different algorithms according to the dataset and based on visualization. We will be able to know how to build a web application using the Flask framework

2.LITERATURE SURVEY

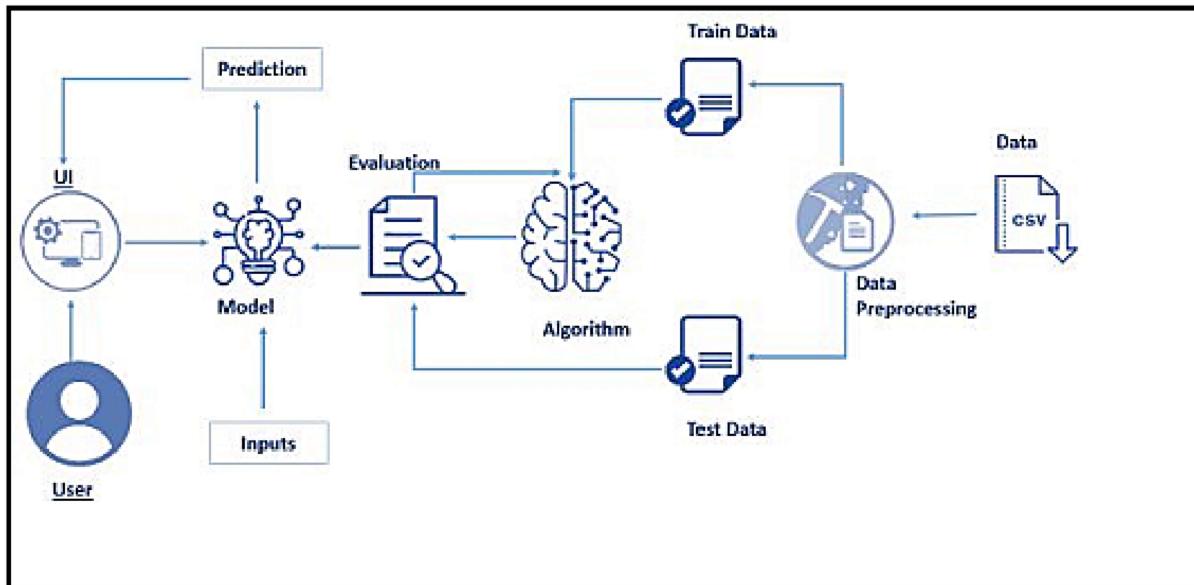
Bag of words statistics and bag of symbols statistics are the typical source information for the toxic comment's detection. Usually, the following statistics-based features are used: length of the comment, number of capital letters, number of exclamation marks, number of question marks, number of spelling errors, number of tokens with non-alphabet symbols, number of abusive, aggressive, and threatening words in the comment, etc.

b. Proposed Solution :

The goal is to create a classifier model that can predict if input text is inappropriate (toxic). Explore the dataset to get a better picture of how the labels are distributed, how they correlate with each other, and what defines toxic or clean comments. Create a baseline score with a simple logistic regression classifier. Explore the effectiveness of multiple machine learning approaches and select the best for this problem. Select the best model and tune the parameters to maximize performance. Build the final model with the best performing algorithm and parameters and test it on a holdout subset of the data.

2. THEORETICAL ANALYSIS

3.1 Block Diagram

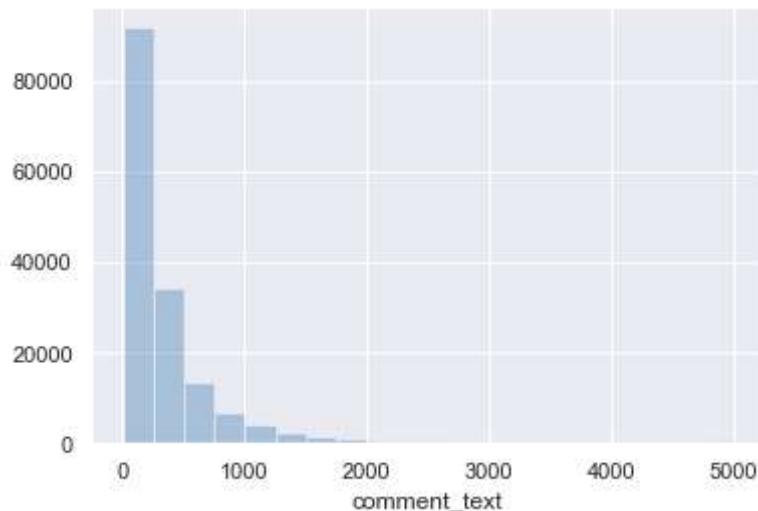


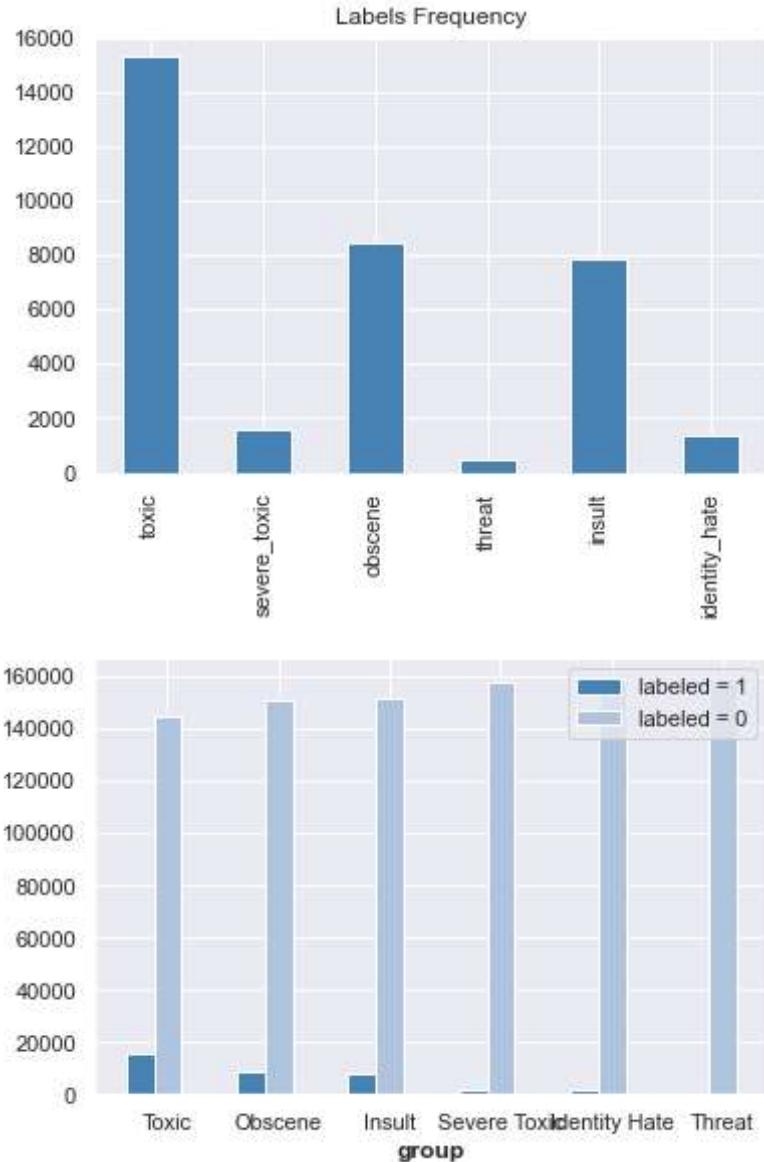
b. Hardware/Software Requirements

1. IDE: Jupyter Notebook, Spyder, Anaconda navigator
2. Programming Language (Back-end): Python 3.7
3. Front-end: HTML, CSS
4. Framework: Flask

4. EXPERIMENTAL INVESTIGATIONS

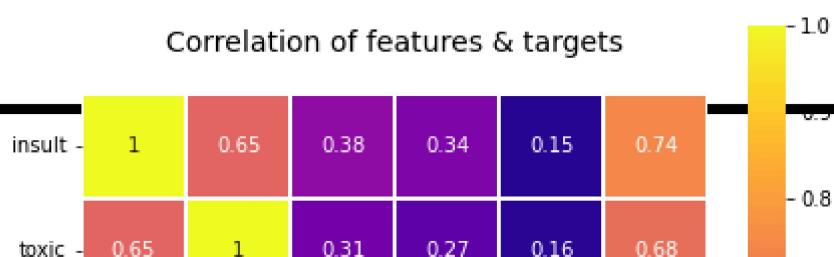
Data Exploration This dataset contains 159,571 comments from Wikipedia. The data consists of one input feature, the string data for the comments, and six labels for different categories of toxic comments: toxic, severe toxic, obscene, threat, insult, and identity hate. The figure on the following page contains a breakdown of how the labels are distributed throughout the dataset, including overlapping data. As you can see in the breakdown, while most comments with other labels are also toxic, not all of them are. Only “severe toxic” is clearly a subcategory of “toxic.” And it’s not close enough to be a labelling error. This suggests that “toxic” is not a catch-all label, but rather a subcategory in itself with a large amount of overlap. Because of this, I’m going to create a seventh label called “any_label” to represent overall toxicity of a comment. From here on in, I’m going to refer to any labelled comments as toxic, and the specific “toxic” label (along with other labels) in quotation marks.



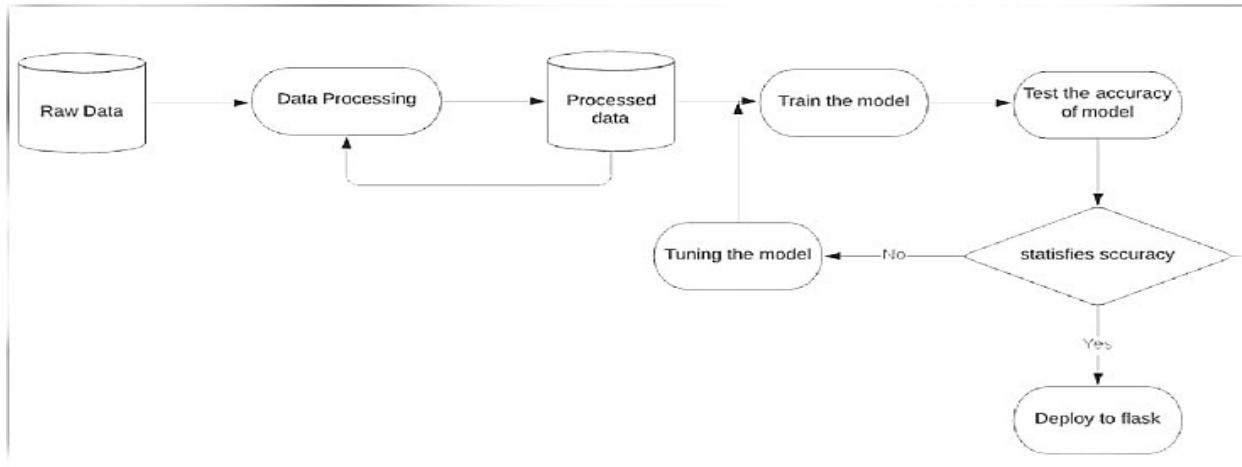


Only 39 percent of the toxic comments have only one label, and the majority have some sort of overlap. I believe that because of this, it will be much more difficult to train a classifier on specific labels than whether they are toxic. This ambiguity and the lack of explanation around it is what led me to select an aggregate label of general toxicity, what I've called “any_label,” as the target. Sixteen thousand two hundred twenty-five out of 159571 comments, or 10.17%, are classified as some category of toxic.

The correlation matrix below provides more insight into these overlapping categories. Threats are not likely to be severely toxic, nor are they likely to be racist or homophobic. But insults are often obscene, and identity hate really doesn't have much overlap at all.



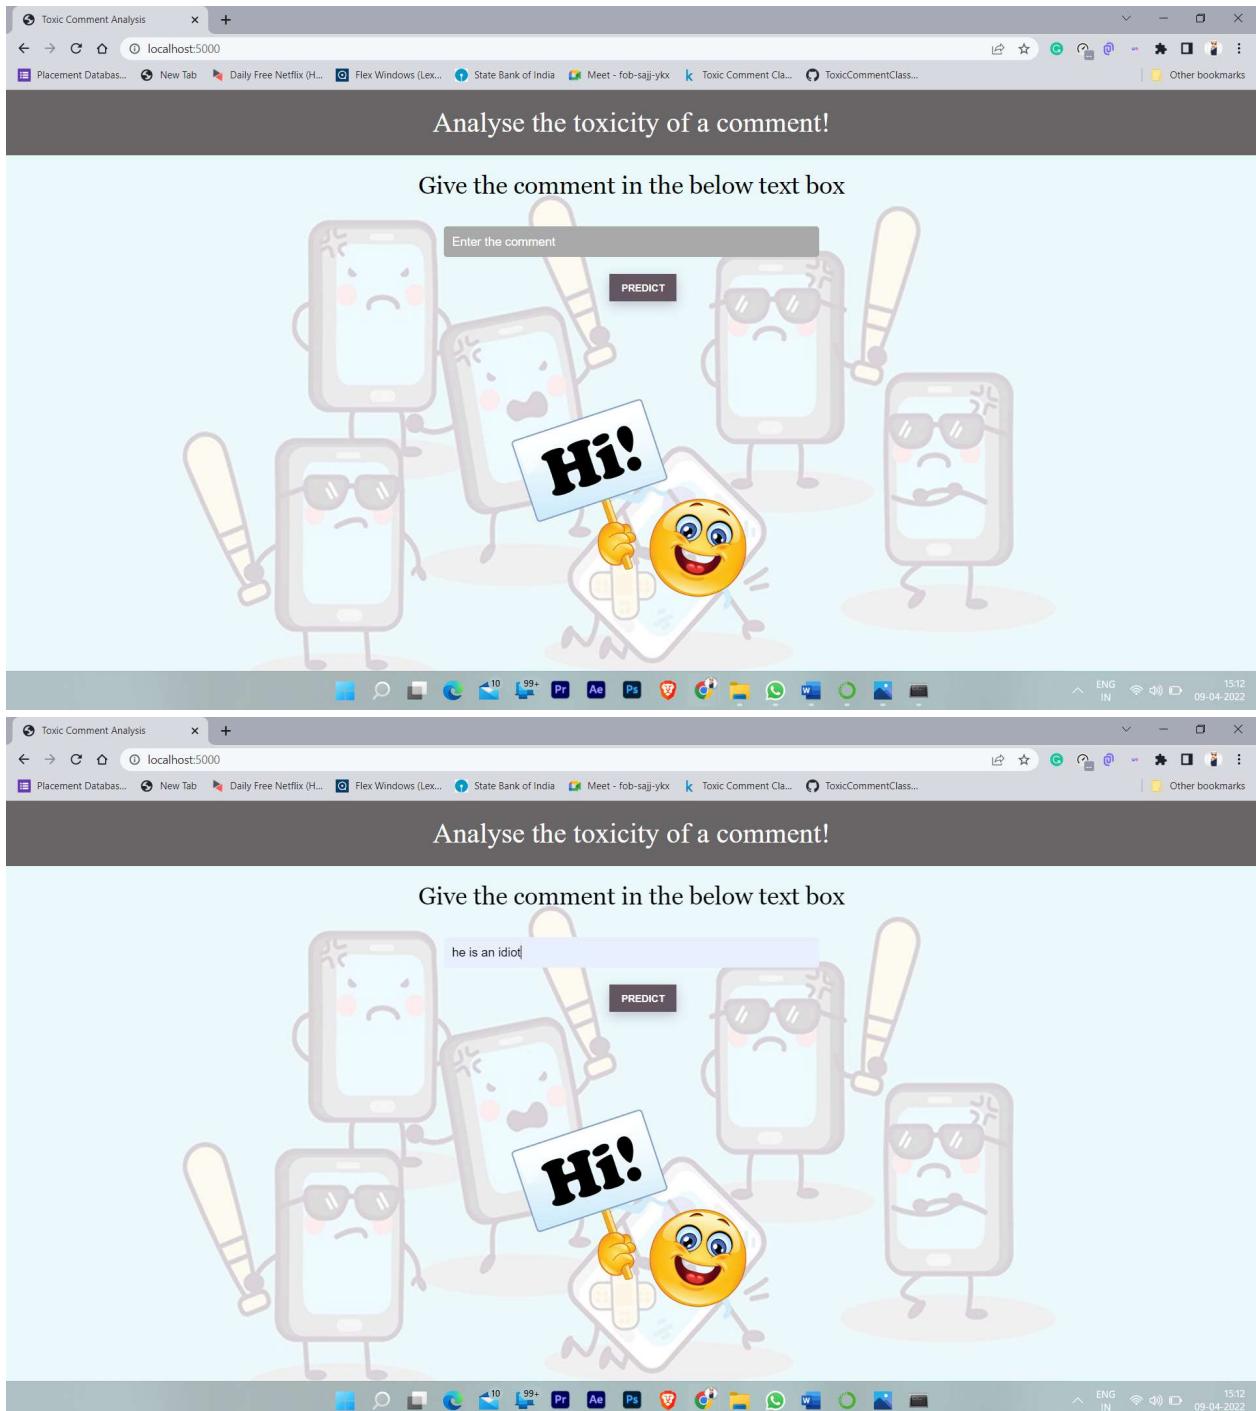
5. FLOW CHART



6. RESULT

We have successfully built the UI interface for classification of toxic comments, in which user

can input any comment and then while pressing predict the user can find the comment is toxic or clean , this UI will help them to find the comment is toxic or clean .



7. ADVANTAGES

1. The interface is user friendly and easy to use and understand even to a person who has very

little/no knowledge .

2. This model uses a huge amount of data the results generated are accurate and reliable.
3. Easy to use & has a user-friendly interface.
4. Results can be improved by training data to our choice of parameter.
5. We can easily do the classification of different comments whether the comment is toxic or clean using natural language processing.

DISADVANTAGES

1. we can't predict the comment if it is different language
2. the modal can't determine properly if the comment is big

8. APPLICATIONS

1. using this modal we can identify the toxic comments
2. Time saving & cost-efficient method as we don t need to read every comment
3. by using this modal we can skip the toxic comments
4. As the same way we can find the person who handles the toxic comments and we can block them
5. using this modal we can clear the unwanted and toxic comments
6. By clearing the waste/toxic data we can reduce the data storage.
7. As this modal is reliable we can use in different sectors to classified the comments.

9. CONCLUSION

A Machine Learning model, has been developed to **classification of toxic**

comments in social networking. A simple, efficient and a versatile model is built keeping in mind the diversity of data, computational complexity and overhead involved in making API calls to the model for prediction. The product can increase the accuracy of finding the toxic comments in social network. It helps to classified the comments and identify the toxic comments.

10.FUTURE SCOPE

1. Recurrent neural networks, despite their increased overhead, could be a very effective solution if GPU resources are available for quick predictions
2. We could use a simple decision tree to feed comments into the model that would be most effective. A few that I can think of are:
 1. Short comments
 2. Long comments
 3. "Hot" threads where the rate of commenting is high and emotions may be high

11. BIBLIOGRAPHY

1. DATASET-- <https://www.kaggle.com/competitions/jigsaw-toxic-comment-classification-challenge/data>
2. Toxic Comment Classification Group Project for MSDS621 Machine Learning at University of San Francisco

12. APPENDIX

a. Source code:

[classifying-multi-label-comments-0-9741-lb.py](#)

S... SI-8... sma... toxic... 2b8... Toxic... Toxic... tian... Toxic... Home... Unt... G... relia... +

localhost:8888/notebooks/classifying-multi-label-comments-0-9741-lb.ipynb#

Placement Databases... New Tab Daily Free Netflix (Lex... Flex Windows (Lex... State Bank of India Meet - fob-saji-ykx Toxic Comment Cla... ToxicCommentClass...

jupyter classifying-multi-label-comments-0-9741-lb Last Checkpoint: 3 minutes ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel) Logout

In [2]: #importing some of the required libraries
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
#get_ipython().run_line_magic('matplotlib', 'inline')
import seaborn as sns
import re
import pickle

In [3]: # We've separate data available for training and testing.
Load training and test data
train_df = pd.read_csv(r'C:\Users\SADHU SUNDHAR\Downloads\jigsaw-toxic-comment-classification-challenge\train.csv\train.csv')
test_df = pd.read_csv(r'C:\Users\SADHU SUNDHAR\Downloads\jigsaw-toxic-comment-classification-challenge\test.csv\test.csv')

DATA DESCRIPTION

In [5]: sns.set(color_codes=True)
comment_len = train_df.comment_text.str.len()
sns.distplot(comment_len, kde=False, bins=20, color="steelblue")

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

Out[5]: <AxesSubplot:xlabel='comment_text'>



S... SI-8... sma... toxic... 2b8... Toxic... Toxic... tian... Toxic... Home... Unt... G... relia... +

localhost:8888/notebooks/classifying-multi-label-comments-0-9741-lb.ipynb#

Placement Databases... New Tab Daily Free Netflix (Lex... Flex Windows (Lex... State Bank of India Meet - fob-saji-ykx Toxic Comment Cla... ToxicCommentClass...

jupyter classifying-multi-label-comments-0-9741-lb Last Checkpoint: 3 minutes ago (autosaved)

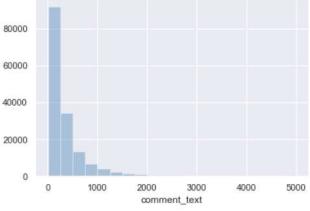
File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel) Logout

DATA DESCRIPTION

In [5]: sns.set(color_codes=True)
comment_len = train_df.comment_text.str.len()
sns.distplot(comment_len, kde=False, bins=20, color="steelblue")

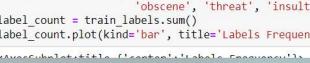
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

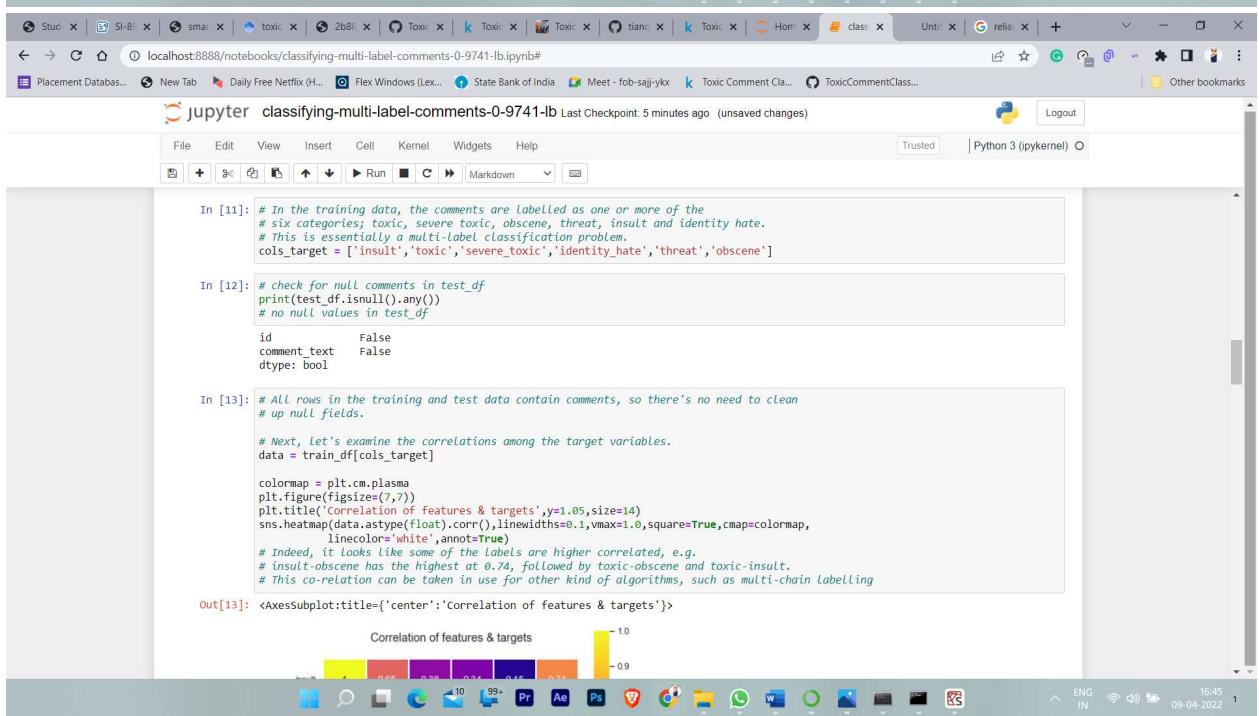
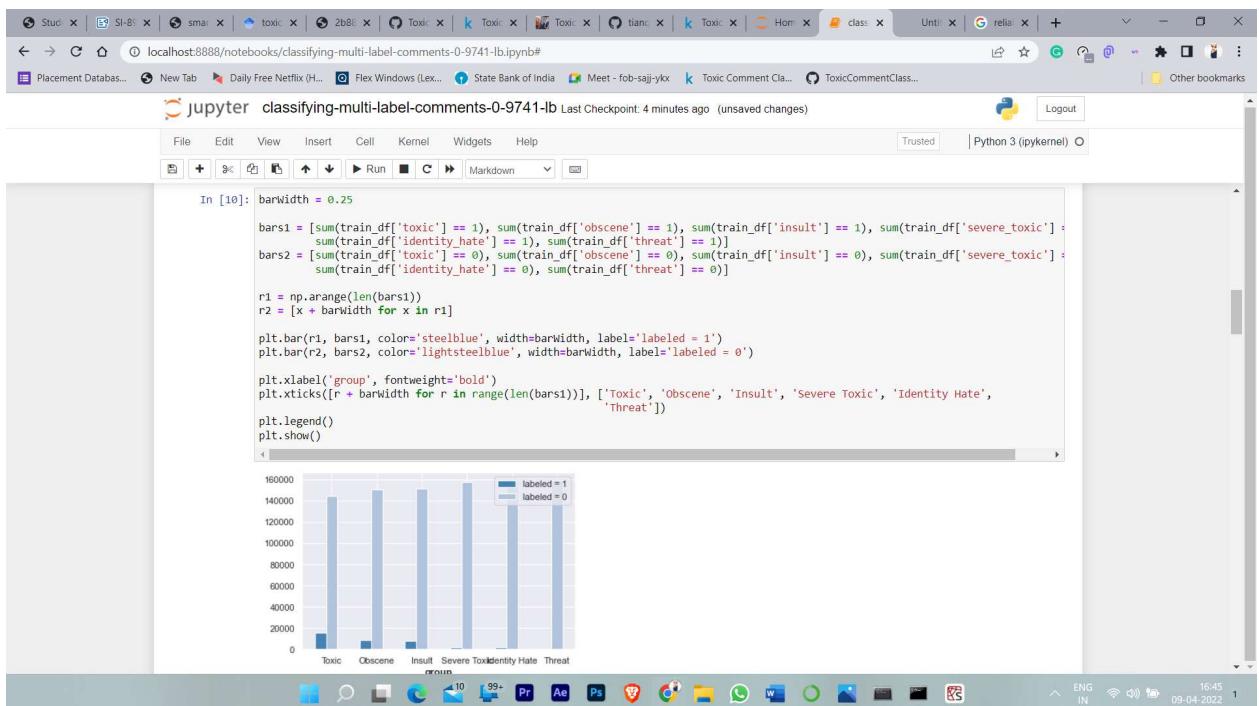
Out[5]: <AxesSubplot:xlabel='comment_text'>

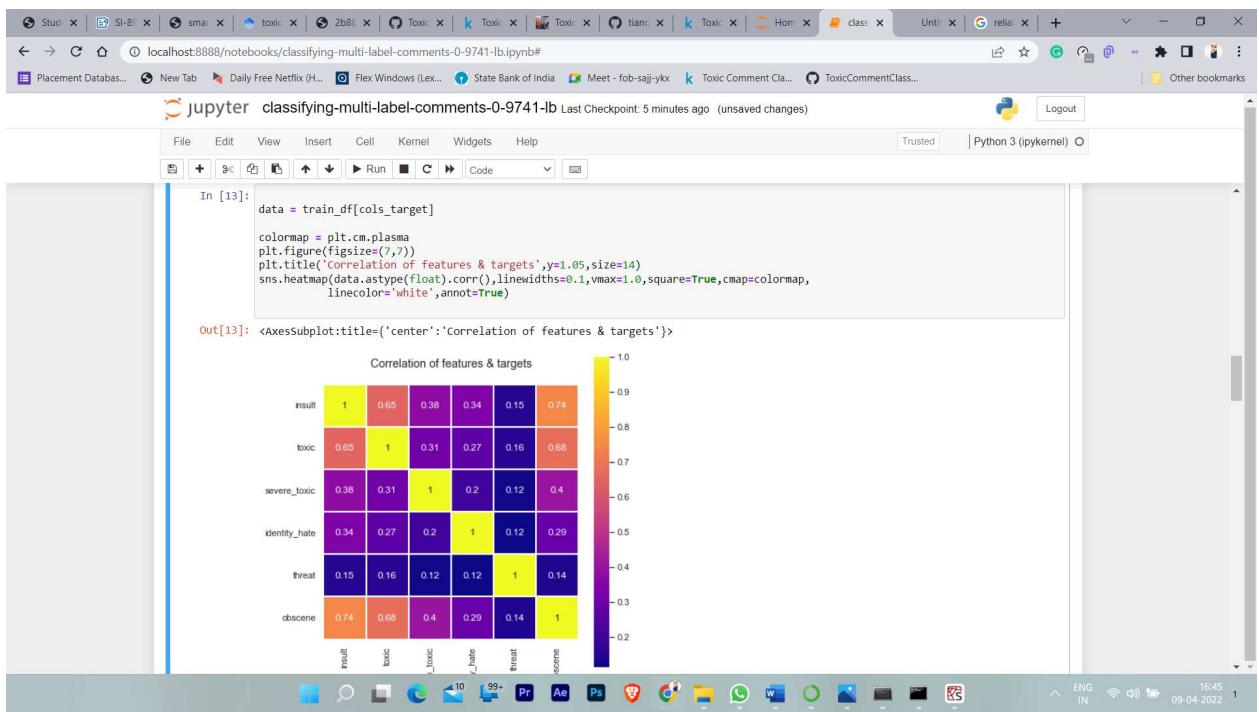


In [7]: train_labels = train_df[['toxic', 'severe_toxic',
'obscene', 'threat', 'insult', 'identity_hate']]
label_count = train_labels.sum()
label_count.plot(kind='bar', title='Labels Frequency', color='steelblue')

Out[7]: <FigureSubplot:title='Labels Frequency'>







Jupyter Notebook: classifying-multi-label-comments-0-9741-lb

DATA Pre-Processing

In [15]:

```
# Define a function to clean up the comment text, basic NLP
def clean_text(text):
    text = text.lower()
    text = re.sub("what's", " what is ", text)
    text = re.sub("n't", " not ", text)
    text = re.sub("ve", " have ", text)
    text = re.sub("can't", " cannot ", text)
    text = re.sub("n't", " not ", text)
    text = re.sub("i'm", " i am ", text)
    text = re.sub("re", " are ", text)
    text = re.sub("d", " would ", text)
    text = re.sub("ll", " will ", text)
    text = re.sub("scuse", " excuse ", text)
    text = re.sub('w', ' ', text)
    text = re.sub('s+', ' ', text)
    text = text.strip(' ')
    return text
```

In [16]:

```
# clean the comment_text in both the datasets.
train_df['comment_text'] = train_df['comment_text'].map(lambda com : clean_text(com))
test_df['comment_text'] = test_df['comment_text'].map(lambda com : clean_text(com))
```

In [17]:

```
# Define all_text from entire train & test data for use in tokenization by Vectorizer
train_text = train_df['comment_text']
test_text = test_df['comment_text']
all_text = pd.concat([train_text, test_text])
```

To [18]:

The screenshot shows a Jupyter Notebook interface running on a Windows desktop. The browser title bar indicates the URL is `localhost:8888/notebooks/classifying-multi-label-comments-0-9741-lb.ipynb#`. The notebook has tabs for various files like 'Study', 'SI-8...', 'sma...', 'toxi...', '2b8...', 'Toxic...', 'Toxic...', 'tian...', 'Toxic...', 'Home...', 'Unit...', 'Grelia...', and 'Logout'. The main content area displays the following Python code:

```
In [18]: #Vectorize the data
# import and instantiate countvectorizer
from sklearn.feature_extraction.text import CountVectorizer
word_vect = CountVectorizer(
    strip_accents='unicode',
    analyzer='word',
    token_pattern='\w{1,}',
    stop_words='english',
    ngram_range=(1, 1))

In [19]: # learn the vocabulary in the training data, then use it to create a document-term matrix
word_vect.fit(all_text)

Out[19]: CountVectorizer(stop_words='english', strip_accents='unicode',
   token_pattern='\w{1,}')

In [20]: # transform the data using the earlier fitted vocabulary, into a document-term matrix
train_features = word_vect.transform(train_text)
test_features = word_vect.transform(test_text)

In [21]: #saving word vectorizer vocab as pkl file to be loaded afterwards
pickle.dump(word_vect.vocabulary_,open('word_feats.pkl','wb'))
```

Below the code, there is explanatory text:

Solving a multi-label classification problem One way to approach a multi-label classification problem is to transform the problem into separate single-class classifier problems. This is known as 'problem transformation'. There are three methods:

Binary Relevance : This is probably the simplest which treats each label as a separate single classification problems. The key assumption here though, is that there are no correlation among the various labels.

Classifier Chains : In this method, the first classifier is trained on the input X. Then the subsequent classifiers are trained on the input X and all previous

This is a second screenshot of the same Jupyter Notebook session. The code and explanatory text are identical to the first screenshot. The browser title bar and the system tray at the bottom right of the screen are also identical, showing the date and time as 09-04-2022 16:46.

The screenshot shows a Jupyter Notebook interface running on a local host. The notebook has three cells:

- In [22]:

```
# import and instantiate the Logistic Regression model
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
logreg = LogisticRegression(C=12.0)
```
- In [23]:

```
# create submission file
submission_binary = pd.read_csv(r"C:\Users\SADHU SUNDHAR\Downloads\jigsaw-toxic-comment-classification-challenge\sample_submission.csv")
```
- In [24]:

```
mapper = {}
for label in cols_target:
    mapper[label] = logreg
    filename = str(label) + '_model.sav'
    print(filename)
    print('... Processing {}'.format(label))
    y = train_df[label]
    # train the model using train_features & y
    mapper[label].fit(train_features, y)

    # saving the fitted model for class "label"
    pickle.dump(mapper[label], open(filename, 'wb'))

# compute the training accuracy
y_pred_X = mapper['insult'].predict(train_features)
print("Training accuracy is {}".format(accuracy_score(y, y_pred_X)))
# compute the predicted probabilities for X test_dtm
test_y_prob = mapper['insult'].predict_proba(test_features)[:,1]
submission_binary['label'] = test_y_prob
```

The output of In [24] shows a ConvergenceWarning from scikit-learn's logistic regression solver.

The screenshot shows the same Jupyter Notebook interface after the execution of In [24]. The output area is filled with repeated error messages from the logistic regression solver, indicating convergence failure due to reaching the iteration limit. The errors are shown in red boxes:

- ... Processing insult
- C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
- Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
- Training accuracy is 0.974437711129215
toxic_model.sav
... Processing toxic
- C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
- Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
- Training accuracy is 0.981381328675435
severe_toxic_model.sav
... Processing severe_toxic
- C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
- Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

The image shows two screenshots of a Jupyter Notebook interface running on a Windows desktop. Both screenshots are identical, displaying a notebook titled "classifying-multi-label-comments-0-9741-lb".

The notebook content consists of several code cells, each processing different categories of toxic comments:

- severe_toxic**:
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1);
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result()
Training accuracy is 0.992160229615657
identity_hate_model.sav
- identity_hate**:
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1);
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result()
Training accuracy is 0.9932819873285246
threat_model.sav
- threat**:
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1);
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result()
Training accuracy is 0.9932819873285246
obscene_model.sav
- obscene**:
Training accuracy is 0.9985085009180866
Training accuracy is 0.984151255553954
- Submission**:
In [33]: # generate submission file
submission_binary.to_csv('submission_binary.csv', index=False)

The notebook interface includes standard Jupyter controls (File, Edit, View, Insert, Cell, Kernel, Widgets, Help) and a Python 3 (ipykernel) kernel selector. The status bar at the bottom of the window shows system information like battery level, signal strength, and date/time.

testing comment toxicity.py

```
File Edit Selection View Go Run Terminal Help testing comment toxicity.py - Visual Studio Code
Get Started testing comment toxicity.py > clean_text
C: > Users > SADHU SUNDHAR > Desktop > sadhu pt2 > testing comment toxicity.py > clean_text
1 import numpy as np
2 import pandas as pd
3 import re
4 from sklearn.feature_extraction.text import CountVectorizer
5 import pickle
6
7 loaded=CountVectorizer(decode_error='replace',vocabulary=pickle.load(open('word_features.pkl','rb')))
8
9 #test_comment='Food was really boring'
10 #test_comment=test_comment.split('delimiter')
11 #result=loaded.transform(test_comment)
12 #print(result)
13
14 #train_df = pd.read_csv('train.csv')
15 #test_df = pd.read_csv('test.csv')
16
17 user_df = pd.DataFrame(columns = ['comment_text'])
18 user_df
19
20 inp_comment = str(input("enter a comment : "))
21 inp_comment
22
23 new_row = {'comment_text':inp_comment}
24 #append a new row to the dataframe formed by user inputs
25 user_df = user_df.append(new_row,ignore_index = True)
26
27
28 def clean_text(text):
29     text = text.lower()
30     text = re.sub(r"What's", "What is ", text)
31     text = re.sub(r"\s+", " ", text)
32     text = re.sub(r've', " have ", text)
33     text = re.sub(r'can't', "cannot ", text)
34     text = re.sub(r'n't', " not ", text)
35     text = re.sub(r'i'm', " i am ", text)
36     text = re.sub(r're', " are ", text)
37     text = re.sub(r'\d', " would ", text)
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
```

Ln 36, Col 42 Spaces: 4 UTF-8 CRLF Python 3.9.7 (base: conda) ⚙ 17:42 09-04-2022

```
File Edit Selection View Go Run Terminal Help testing comment toxicity.py - Visual Studio Code
Get Started testing comment toxicity.py > clean_text
C: > Users > SADHU SUNDHAR > Desktop > sadhu pt2 > testing comment toxicity.py > clean_text
37 text = re.sub(r"\d", " would ", text)
38 text = re.sub(r"\ll", " will ", text)
39 text = re.sub(r"\scuse", " excuse ", text)
40 text = re.sub('W', ' ', text)
41 text = re.sub('s+', ' ', text)
42 text = text.strip(' ')
43 return text
44
45 # clean the comment_text in train_df [Thanks to Pulkit Jha for the useful pointer.]
46 #train_df['comment_text'] = train_df['comment_text'].map(lambda com : clean_text(com))
47
48 # clean the comment_text in test_df [Thanks, Pulkit Jha.]
49 #test_df['comment_text'] = test_df['comment_text'].map(lambda com : clean_text(com))
50
51 # clean the comment_text in user_df [Thanks, Pulkit Jha.]
52 user_df['comment_text'] = user_df['comment_text'].map(lambda com : clean_text(com))
53 user_df
54
55
56 #train_text = train_df['comment_text']
57 #test_text = test_df['comment_text']
58 user_text = user_df['comment_text']
59
60
61 #train_features = loaded.transform(train_text)
62 # transform the test data using the earlier fitted vocabulary, into a document-term matrix
63 #test_features = loaded.transform(test_text)
64
65 user_features = loaded.transform(user_text)
66 print(user_features)
67
68 cols_target = ['obscene','insult','toxic','severe_toxic','identity_hate','threat']
69
70 from sklearn.linear_model import LogisticRegression
71 from sklearn.metrics import accuracy_score
72
73 # create submission file
```

Ln 36, Col 42 Spaces: 4 UTF-8 CRLF Python 3.9.7 (base: conda) ⚙ 17:42 09-04-2022

```
C:\ > SADHU SUNDHAR > Desktop > sadhu pt2 > testing comment toxicity.py > clean_text  
#SUBMISSION_BINARY.columns  
76  
77 lst= []  
78 mapper = {}  
79 for label in cols_target:  
80     filename = str(label+'_.model.sav')  
81     filename  
82     model = pickle.load(open(filename, 'rb'))  
83  
84     print('... Processing {}'.format(label))  
85     user_y_prob = model.predict_proba(user_features)[:,1]  
86     print(label,"-",user_y_prob[0])  
87     lst.append([label,user_y_prob])  
88     #submission_binary[label] = test_y_prob  
89  
90 print(lst)  
91 #['obscene', array([0.05674509]),  
92 # ['insult', array([0.07231278])], ['toxic', array([0.10434529])], ['severe_toxic', array([0.01385649])], ['identity_hate', array([0.01463043])], ['threat', array([0.00486619])]  
93 final=[]  
94 for i in lst:  
95     if i[1]>0.5:  
96         final.append(i[0])  
97  
98 if not len(final):  
99     text = "The comment is clean"  
100 else:  
101     text= "the comment is toxic"  
102  
103 print(text)  
104  
105  
106  
107  
108  
109  
110  
111  
112  
Ln 36, Col 42 Spaces: 4 UTF-8 CRLF Python 3.9.7 (base: conda) ⚙ 🔍 🔍
```



```
C:\ > Users > SADHU SUNDHAR > Desktop > sadhu pt2 > testing comment toxicity.py > ...  
1 import numpy as np  
2 import pandas as pd  
3 import re  
4 from sklearn.feature_extraction.text import CountVectorizer  
5 import pickle  
6  
7 loaded=CountVectorizer(decode_error='replace',vocabulary=pickle.load(open('word_features.pkl','rb')))  
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL  
PS C:\Users\SADHU SUNDHAR> & C:/ProgramData/Anaconda3/python.exe "c:/Users/SADHU SUNDHAR/Desktop/sadhu pt2/testing comment toxicity.py"  
enter a comment : he is an idiot  
(0, 140485) 1  
... Processing obscene  
obscene : 0.20713966072594417  
... Processing insult  
insult : 0.7006629806577194  
... Processing toxic  
toxic : 0.9858088272304611  
... Processing severe_toxic  
severe_toxic : 0.0213241557866  
... Processing identity_hate  
identity_hate : 0.013324584695474614  
... Processing threat  
threat : 0.0017126708631557866  
[['obscene', array([0.20713966])], ['insult', array([0.70066298])], ['toxic', array([0.98580883])], ['severe_toxic', array([0.02132415])], ['identity_hate', array([0.01332458])], ['threat', array([0.00171207])]]  
the comment is toxic  
PS C:\Users\SADHU SUNDHAR> 🔍 🔍
```


Ln 4, Col 60 Spaces: 4 UTF-8 CRLF Python 3.9.7 (base: conda) ⚙ 🔍 🔍

commentApp.py

```
File Edit Selection View Go Run Terminal Help
Get Started testing comment toxicity.py commentApp.py
C: > Users > SADHU SUNDHAR > Desktop > pt.2 > Commit-Toxicity-Multi-Class-Classification-main > Commit-Toxicity-Multi-Class-Classification-main > Flask > commentApp.py ...
1 import numpy as np
2 import pandas as pd
3 import re
4 from sklearn.feature_extraction.text import CountVectorizer
5 import pickle
6 from sklearn.linear_model import LogisticRegression
7 from sklearn.metrics import accuracy_score
8 from flask import Flask, request, jsonify, render_template, url_for
9
10 loaded=CountVectorizer(decode_error='replace',vocabulary=pickle.load(open('word_features.pkl','rb')))
11
12 app = Flask(__name__)
13
14 def clean_text(text):
15     text = text.lower()
16     text = re.sub(r"What's", " what is ", text)
17     text = re.sub(r"\ s", " ", text)
18     text = re.sub(r"\ ve", " have ", text)
19     text = re.sub(r"can't", " cannot ", text)
20     text = re.sub(r"n't", " not ", text)
21     text = re.sub(r"i'm", " i am ", text)
22     text = re.sub(r"re", " are ", text)
23     text = re.sub(r"\d", " would ", text)
24     text = re.sub(r"\ll", " will ", text)
25     text = re.sub(r"\scuse", " excuse ", text)
26     text = re.sub(r'\W', ' ', text)
27     text = re.sub(r'\s+', ' ', text)
28     text = text.strip(' ')
29     return text
30
31 @app.route('/')
32 def landingpage():
33     img_url = url_for('static', filename = 'images/hello.png')
34     print(img_url)
35     flag=0
36     return render_template('toxic.html', flag=flag)
37
```

Ln 1, Col 1 Spaces: 4 UTF-8 CRLF Python 3.9.7 (base: conda) 17:52 09-04-2022

```
File Edit Selection View Go Run Terminal Help
Get Started testing comment toxicity.py commentApp.py
C: > Users > SADHU SUNDHAR > Desktop > pt.2 > Commit-Toxicity-Multi-Class-Classification-main > Commit-Toxicity-Multi-Class-Classification-main > Flask > commentApp.py ...
37
38 @app.route('/predict')
39 @app.route('/', methods = ['GET','POST'])
40 def predict():
41     if request.method == 'GET':
42         img_url = url_for('static',filename = 'images/hello.png')
43         return render_template('toxic.html', url=img_url)
44     if request.method == 'POST':
45         comment = request.form['comment']
46         new_row = {'comment_text':comment}
47         user_df = pd.DataFrame(columns = ['comment_text'])
48         user_df = user_df.append(new_row,ignore_index = True)
49         user_df['comment_text'] = user_df['comment_text'].map(lambda com : clean_text(com))
50         user_text = user_df['comment_text']
51         user_features = loaded.transform(user_text)
52         cols_target = ['obscene', 'insult', 'toxic', 'severe_toxic', 'identity_hate', 'threat']
53         lst= []
54         mapper = {}
55         for label in cols_target:
56             filename = str(label+'_model.sav')
57             print(filename)
58             model = pickle.load(open(filename, 'rb'))
59             print('... Processing {}'.format(label))
60             user_y_prob = model.predict_proba(user_features)[:,1]
61             print(label,":",user_y_prob[0])
62             lst.append([label,user_y_prob])
63         print(lst)
64
65         final=[]
66         flag=0
67         for i in lst:
68             if i[1]>0.5:
69                 final.append(i[0])
70                 flag=2
71         if not len(final):
72             text = "Yayy!! The comment is clean"
73             img_url = url_for('static',filename = 'images/happy.png')
```

Ln 1, Col 1 Spaces: 4 UTF-8 CRLF Python 3.9.7 (base: conda) 17:52 09-04-2022

```
53     lst= []
54     mapper = {}
55     for label in cols_target:
56         filename = str(label+'_.model.sav')
57         print(filename)
58         model = pickle.load(open(filename, 'rb'))
59         print('... Processing {}'.format(label))
60         user_y_prob = model.predict_proba(user_features)[:,1]
61         print(label,":",user_y_prob[0])
62         lst.append([label,user_y_prob])
63     print(lst)
64
65     final=[]
66     flag=0
67     for i in lst:
68         if i[1]>0.5:
69             final.append(i[0])
70             flag+=1
71     if not len(final):
72         text = "Yayy!! The comment is clean"
73         img_url = url_for('static',filename = 'images/happy.png')
74         flag=1
75         print(img_url)
76     else:
77         text="The comment is "
78         for i in final:
79             text = text+i+
80             img_url = url_for('static',filename = 'images/toxic.png')
81         print(text)
82     return render_template('toxic.html',ypred = text,url= img_url,flag=flag)
83
84 if __name__ == '__main__':
85     app.run(host = 'localhost', debug = False , threaded = False)
86
87
```

Ln 1, Col 1 Spaces: 4 UTF-8 CRLF Python 3.9.7 (base: conda) 17:52 09-04-2022 1 ENG IN

```
53     lst= []
54     mapper = {}
55     for label in cols_target:
56         filename = str(label+'_.model.sav')
57         print(filename)
58         model = pickle.load(open(filename, 'rb'))
59         print('... Processing {}'.format(label))

```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL

```
fication-main/Flask/commentApp.py"
* Serving Flask app "commentApp" (lazy loading)
* Environment: production
WARNING: This is a development server. Do not use it in a production deployment.
Use a production WSGI server instead.
* Debug mode: off
* Running on http://localhost:5000/ (Press CTRL+C to quit)
```

Python Python

Ln 1, Col 1 Spaces: 4 UTF-8 CRLF Python 3.9.7 (base: conda) 17:52 09-04-2022 1 ENG IN

b. UI Output Screenshot

FIG:1 HOME PAGE

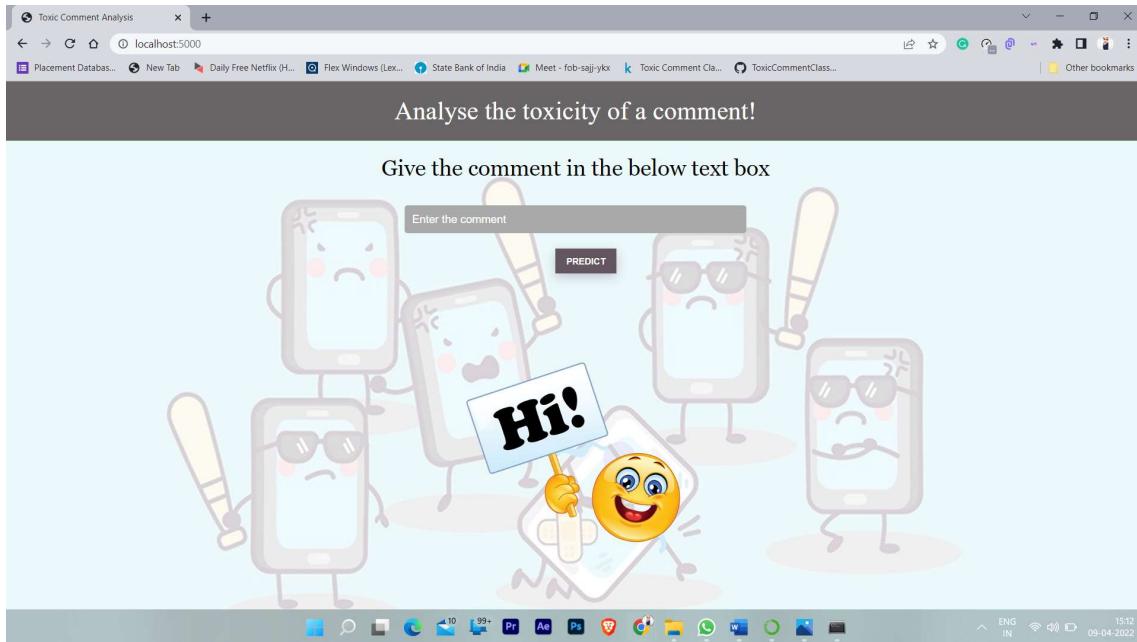


FIG:2 ENTER A TOXIC COMMENT

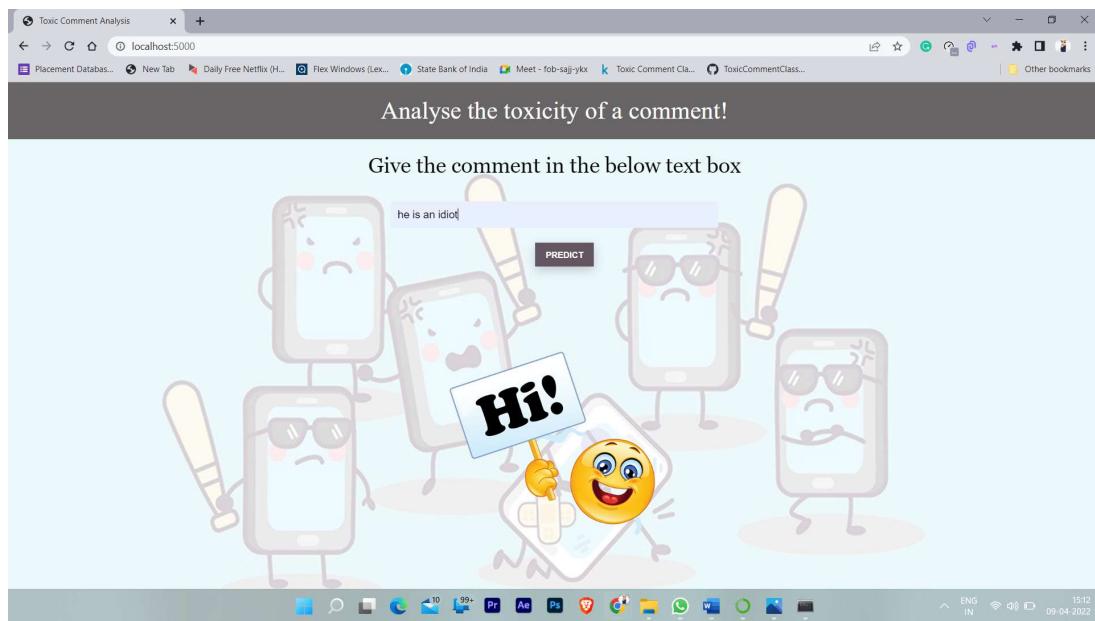


FIG:3 SHOWING OUTPUT THE COMMENT IS TOXIC

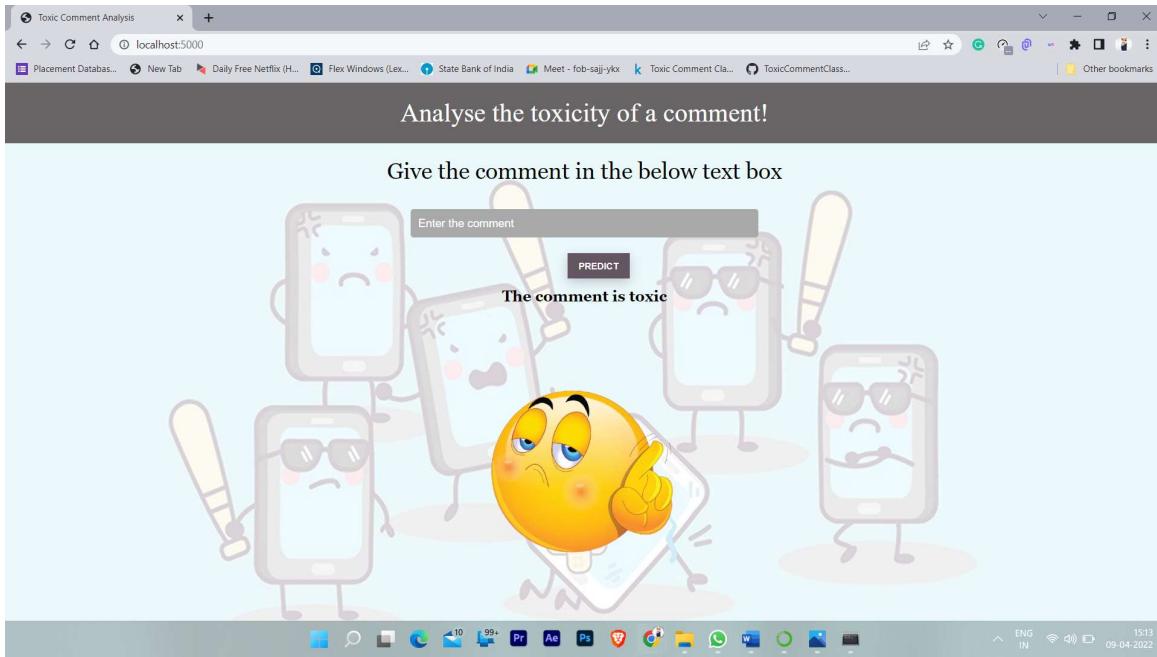


FIG:4 ENTERING NORMAL COMMENT

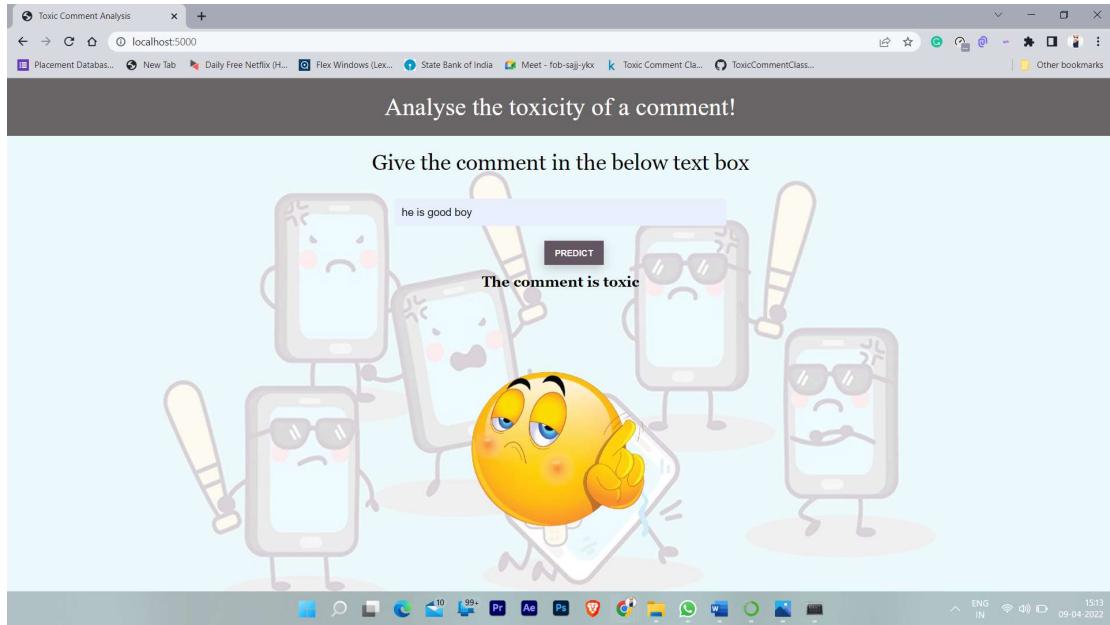
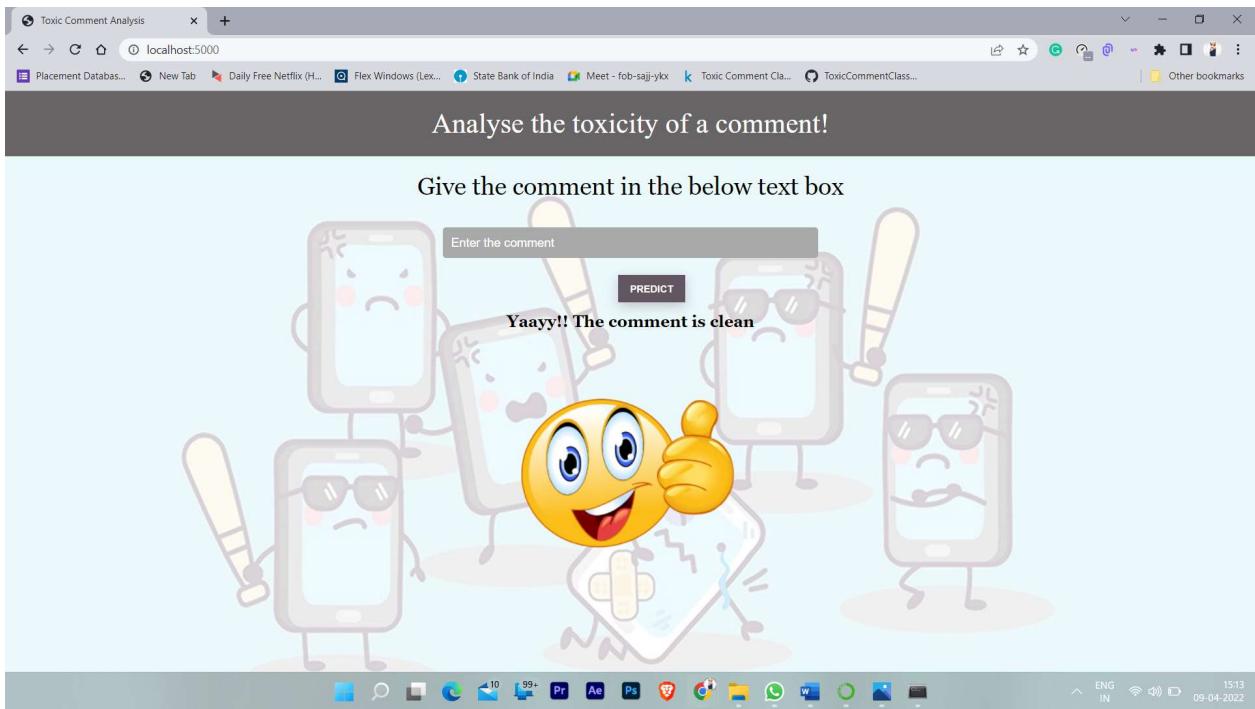


FIG:3 SHOWING OUTPUT THE COMMENT IS CLEAN



THANK YOU

Submitted By

PENTAPATI ESWAR AKHIL

39110765